

Personal Firewalls - An Introduction to Firewall Administration

David Tran - A00801942

COMP 6D

COMP 8006

British Columbia Institute of Technology

Aman Abdulla

Thursday, February 6 2014

Table of Contents

BACKGROUND	3
TOOLS & EQUIPMENT	3
Hardware	3
Software	3
TESTING PROCEDURE	3
Test Cases Table	3
Test Case Evidence & Details	5
CONCLUSION	19
APPENDIX	20

Background

The purpose of this assignment is to familiarize the budding administrator with the syntax and implementation of firewall rules on a Linux machine. The objective is to design a set of rules that follow the criteria of Assignment 1, as well as documenting a set of test cases properly so that the routine becomes natural. Testing in this assignment will require a tertiary computer terminal located in the same local area network. In this test case, we will assume that the tertiary computer will not have anything but the default firewall in place.

Tools & Equipment

Hardware

- 8GB RAM
- Intel i5 Quad Core
- 500GB HDD
- two computer terminals

Software

- Fedora Linux 19 64-bit
- hping3
- Shell Script
- Wireshark
- iptables
- ssh
- dhclient

Testing Procedure

Test Cases Table

Case #	Test Case	Tools Used	Expected Outcome	Results
1	Verify that iptables settings are changed after script is ran	iptables -L -n -v -x	Rules are changed according to script	PASSED. See details.
2	DNS packets are allowed to and from the host	Wireshark, Google Chrome	Packets are allowed to travel to and from the machine	PASSED. See details.
3	DHCP packets are allowed to and from the host	ifconfig, dhclient	Packets are allowed to travel to and from the machine	PASSED. See details.
4	SSH traffic is allowed and accounted for	ssh, iptables -L -n -v -x	SSH traffic is permitted inbound	PASSED. See

			and outbound	details.
5	WWW(80) traffic is allowed and accounted for	hping3, iptables -L -n -v -x	WWW traffic is permitted inbound and outbound	PASSED. See details.
6	WWW(443) traffic is allowed and accounted for	hping3, iptables -L -n -v -x	WWW traffic is permitted inbound and outbound	PASSED. See details.
7	WWW(80) traffic is disallowed and accounted for when source ports are less than 1024	hping3, iptables -L -n -v -x	WWW traffic is not permitted inbound and outbound using port 0	PASSED. See details.
8	WWW(80) traffic is disallowed and accounted for when source ports are less than 1024	hping3, iptables -L -n -v -x	WWW traffic is not permitted inbound and outbound using port 1023	PASSED. See details.
9	WWW(80) traffic is disallowed and accounted for when source ports are less than 1024	hping3, iptables -L -n -v -x	WWW traffic is permitted inbound and outbound using port 1024	PASSED. See details.
10	WWW(443) traffic is disallowed and accounted for when source ports are less than 1024	hping3, iptables -L -n -v -x	WWW traffic is not permitted inbound and outbound using port 0	PASSED. See details.
11	WWW(443) traffic is disallowed and accounted for when source ports are less than 1024	hping3, iptables -L -n -v -x	WWW traffic is not permitted inbound and outbound using port 1023	PASSED. See details.
12	WWW(443) traffic is disallowed and accounted for when source ports are less than 1024	hping3, iptables -L -n -v -x	WWW traffic is permitted inbound and outbound using port 1024	PASSED. See details.
13	Incoming packets to port 0 are dropped (TCP)	hping3, iptables -L -n -v -x,	Incoming packets are dropped	PASSED. See details.
14	Incoming packets to port 0 are dropped (UDP)	hping3, iptables -L -n -v -x,	Incoming packets are dropped	PASSED. See details.
15	WWW(80) incoming traffic is accounted for	hping3, iptables -L -n -v -x	Instances of www-traffic is incremented	PASSED. See details.

16	WWW(443) incoming traffic is accounted for	hping3, iptables -L -n -v -x	Instances of www-traffic is incremented	PASSED. See details.
17	WWW(80) outgoing traffic is accounted for	hping3, iptables -L -n -v -x	Instances of www-traffic is incremented	PASSED. See details.
18	WWW(443) outgoing traffic is accounted for	hping3, iptables -L -n -v -x	Instances of www-traffic is incremented	PASSED. See details.
19	SSH(22) incoming traffic is accounted for	ssh, hping3, iptables -L -n -v -x	Instances of ssh-traffic is incremented	PASSED. See details.
20	SSH(22) outgoing traffic is accounted for	ssh, hping3, iptables -L -n -v -x	Instances of ssh-traffic is incremented	PASSED. See details.
21	All inbound SYN packets are dropped	hping3	SYN packets are dropped; exceptions are permitted	PASSED. See details.
22	All other undefined ports with inbound traffic are accounted for (TCP)	hping3, iptables -L -n -v -x	Ports 22, 80, 443 will not increment, all others will increment	PASSED. See details.
23	All other undefined ports with outbound traffic are accounted for (TCP)	hping3, iptables -L -n -v -x	Ports 22, 80, 443 will not increment, all others will increment	PASSED. See details.
24	All other undefined ports with inbound traffic are accounted for (UDP)	hping3, iptables -L -n -v -x	Ports 22, 80, 443 will not increment, all others will increment	PASSED. See details.
25	All other undefined ports with inbound traffic are accounted for (UDP)	hping3, iptables -L -n -v -x	Ports 22, 80, 443 will not increment, all others will increment	PASSED. See details.

Test Case Evidence & Details

(1) Verify that iptables settings are changed after script is ran

Before: iptables -L -n -v -x

```
[root@DataComm ~]# iptables -L -n -v -x
Chain INPUT (policy DROP 0 packets, 0 bytes)
  pkts      bytes target     prot opt in     out     source            destination
   796    104249 ACCEPT      all  --  *      *        0.0.0.0/0         0.0.0.0/0

Chain FORWARD (policy DROP 0 packets, 0 bytes)
  pkts      bytes target     prot opt in     out     source            destination
    0         0 ACCEPT      all  --  *      *        0.0.0.0/0         0.0.0.0/0

Chain OUTPUT (policy DROP 0 packets, 0 bytes)
  pkts      bytes target     prot opt in     out     source            destination
   718    237901 ACCEPT      all  --  *      *        0.0.0.0/0         0.0.0.0/0
[root@DataComm ~]#
```

After: ./firewall_1.sh

```
[root@DataComm ~]# ./firewall_1.sh
iptables: Saving firewall rules to /etc/sysconfig/iptables:[ OK ]
Redirecting to /bin/systemctl restart iptables.service
Chain INPUT (policy DROP 0 packets, 0 bytes)
  pkts      bytes target     prot opt in     out     source            destination
    0         0 tcp_in_chain tcp -- em1    *        0.0.0.0/0         0.0.0.0/0
    0         0 udp_in_chain udp -- em1    *        0.0.0.0/0         0.0.0.0/0
    0         0 DROP      tcp -- em1    *        0.0.0.0/0         192.168.0.2      tcp dpt:0
    0         0 DROP      udp -- em1    *        0.0.0.0/0         192.168.0.2      udp dpt:0
    0         0 DROP      tcp -- em1    *        0.0.0.0/0         0.0.0.0/0         tcp flags:!0x17/0x02 state NEW

Chain FORWARD (policy DROP 0 packets, 0 bytes)
  pkts      bytes target     prot opt in     out     source            destination

Chain OUTPUT (policy DROP 0 packets, 0 bytes)
  pkts      bytes target     prot opt in     out     source            destination
    0         0 tcp_out_chain tcp -- *      em1    0.0.0.0/0         0.0.0.0/0
    0         0 udp_out_chain udp -- *      em1    0.0.0.0/0         0.0.0.0/0
    0         0 DROP      tcp -- *      em1    192.168.0.2       0.0.0.0/0         tcp spt:0
    0         0 DROP      udp -- *      em1    192.168.0.2       0.0.0.0/0         udp spt:0

Chain noness-traffic (10 references)
  pkts      bytes target     prot opt in     out     source            destination
    0         0 ACCEPT      all  --  *      *        0.0.0.0/0         0.0.0.0/0
    0         0 ACCEPT      udp  --  *      *        0.0.0.0/0         0.0.0.0/0         udp dpt:53
    0         0 ACCEPT      udp  --  *      *        0.0.0.0/0         0.0.0.0/0         udp spt:53
    0         0 ACCEPT      tcp  --  *      *        0.0.0.0/0         0.0.0.0/0         tcp dpt:53
    0         0 ACCEPT      tcp  --  *      *        0.0.0.0/0         0.0.0.0/0         tcp spt:53
    0         0 ACCEPT      udp  --  *      em1    0.0.0.0/0         0.0.0.0/0         udp dpts:68:67
    0         0 ACCEPT      udp  --  em1    *        0.0.0.0/0         0.0.0.0/0         udp spts:67:68

Chain ssh-traffic (5 references)
  pkts      bytes target     prot opt in     out     source            destination
    0         0 ACCEPT      all  --  *      *        0.0.0.0/0         0.0.0.0/0
    0         0 ACCEPT      tcp  --  *      *        0.0.0.0/0         192.168.0.2       tcp dpt:22 state NEW,ESTABLISHED
    0         0 ACCEPT      tcp  --  *      *        192.168.0.2       0.0.0.0/0         tcp spt:22 state ESTABLISHED
    0         0 ACCEPT      tcp  --  *      *        192.168.0.2       0.0.0.0/0         tcp dpt:22 state NEW,ESTABLISHED
    0         0 ACCEPT      tcp  --  *      *        0.0.0.0/0         192.168.0.2       tcp spt:22 state ESTABLISHED

Chain tcp_in_chain (1 references)
  pkts      bytes target     prot opt in     out     source            destination
    0         0 ACCEPT      all  --  *      *        0.0.0.0/0         0.0.0.0/0
    0         0 ssh-traffic tcp  --  *      *        0.0.0.0/0         0.0.0.0/0         tcp spt:22
    0         0 www-traffic tcp  --  *      *        0.0.0.0/0         0.0.0.0/0         multiport sports 80,443
    0         0 noness-traffic all -- *      *        0.0.0.0/0         0.0.0.0/0
```

(2) DNS Packets are allowed to and from the host

To begin, we opened a browser and navigated to Google.ca (142.232.191.38)


Inbound DNS Packets:

Filter: dns && ip.dst eq 192.168.0.2

Expression... Clear Apply Save

No.	Time	Source	Destination	Protoccl	Length	Info
80	30.775388000	142.232.191.38	192.168.0.2	DNS	121	Standard query response 0x22b2
183	32.199753000	142.232.191.38	192.168.0.2	DNS	117	Standard query response 0x9441
193	32.221983000	142.232.191.38	192.168.0.2	DNS	135	Standard query response 0xab1b
208	32.240683000	142.232.191.38	192.168.0.2	DNS	272	Standard query response 0xfd37
209	32.240694000	142.232.191.38	192.168.0.2	DNS	149	Standard query response 0xeb65
211	32.240784000	142.232.191.38	192.168.0.2	DNS	162	Standard query response 0xd6be
214	32.240871000	142.232.191.38	192.168.0.2	DNS	251	Standard query response 0xfd59
215	32.240881000	142.232.191.38	192.168.0.2	DNS	91	Standard query response 0x13a9
217	32.240943000	142.232.191.38	192.168.0.2	DNS	154	Standard query response 0x950e
219	32.242548000	142.232.191.38	192.168.0.2	DNS	91	Standard query response 0x169c

Outbound DNS Packets:

Filter:		dns && ip.src eq 192.168.0.2			Expression...	Clear	Apply	Save
No.	Time	Source	Destination	Protoccl	Length	Info		
79	30.774443000	192.168.0.2	142.232.191.38	DNS	73	Standard query 0x22b2	A	www.google.ca
173	32.181664000	192.168.0.2	142.232.191.38	DNS	69	Standard query 0x9441	A	google.ca
191	32.221147000	192.168.0.2	142.232.191.38	DNS	78	Standard query 0xab1b	A	accounts.google.ca
201	32.238527000	192.168.0.2	142.232.191.38	DNS	75	Standard query 0xfd37	A	apis.google.com
202	32.238579000	192.168.0.2	142.232.191.38	DNS	72	Standard query 0xeb65	A	id.google.ca
203	32.238660000	192.168.0.2	142.232.191.38	DNS	85	Standard query 0xd6be	A	lh4.googleusercontent.com
204	32.238739000	192.168.0.2	142.232.191.38	DNS	75	Standard query 0xfd59	A	plus.google.com
205	32.238785000	192.168.0.2	142.232.191.38	DNS	75	Standard query 0x13a9	A	ssl.gstatic.com
206	32.238876000	192.168.0.2	142.232.191.38	DNS	74	Standard query 0x950e	A	www.google.com
210	32.240752000	192.168.0.2	142.232.191.38	DNS	75	Standard query 0x169c	A	www.gstatic.com

(3) DHCP Packets are allowed to and from the host

Firstly, in Terminal, we deactivated our network card and then brought it back up again using:
 ipconfig em1 down; ipconfig em1 up;

```
[root@DataComm ~]# ifconfig em1 down
[root@DataComm ~]# ifconfig em1 up
[root@DataComm ~]# ifconfig
em1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::7a2b:cbff:fea3:4325 prefixlen 64 scopeid 0x20<link>
    ether 78:2b:cb:a3:43:25 txqueuelen 1000 (Ethernet)
    RX packets 152599 bytes 71349241 (68.0 MiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 62561 bytes 14364078 (13.6 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
    device interrupt 20 memory 0xela00000-ela20000

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 0 (Local Loopback)
    RX packets 284 bytes 22708 (22.1 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 284 bytes 22708 (22.1 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

Afterwards, we ran this command:
 dhclient em1

```

[root@DataComm ~]# dhclient em1
[root@DataComm ~]# ifconfig
em1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.0.2 netmask 255.255.255.0 broadcast 192.168.0.255
    inet6 fe80::7a2b:cbff:fea3:4325 prefixlen 64 scopeid 0x20<link>
    ether 78:2b:cb:a3:43:25 txqueuelen 1000 (Ethernet)
    RX packets 152806 bytes 71443815 (68.1 MiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 62757 bytes 14394570 (13.7 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
    device interrupt 20 memory 0xela00000-ela20000

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 0 (Local Loopback)
    RX packets 284 bytes 22708 (22.1 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 284 bytes 22708 (22.1 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

```

We were then able to see that we were dedicated the same IPv4 Address again:
192.168.0.2

(4) SSH traffic is allowed to and from the host

In a second computer terminal, assume that this terminal is the client. We entered the following command into Terminal:

ssh 192.168.0.2

Here's the Wireshark capture of inbound and outbound traffic for SSHv2

Filter: ip.src eq 192.168.0.11 ip.dst eq 192.168.0.11 && s Expression... Clear Apply Save						
No.	Time	Source	Destination	Protoccl	Length	Info
18	8.329554000	192.168.0.11	192.168.0.2	SSHv2	87	Encrypted request packet len=21
20	8.335308000	192.168.0.2	192.168.0.11	SSHv2	87	Encrypted response packet len=21
22	8.336282000	192.168.0.11	192.168.0.2	SSHv2	1898	Client: Key Exchange Init
25	8.336415000	192.168.0.2	192.168.0.11	SSHv2	162	Server: Key Exchange Init
27	8.338308000	192.168.0.11	192.168.0.2	SSHv2	146	Client: Diffie-Hellman Key Exchan
28	8.341110000	192.168.0.2	192.168.0.11	SSHv2	378	Server: New Keys
29	8.343920000	192.168.0.11	192.168.0.2	SSHv2	82	Client: New Keys
31	8.383386000	192.168.0.11	192.168.0.2	SSHv2	118	Encrypted request packet len=52
33	8.383502000	192.168.0.2	192.168.0.11	SSHv2	118	Encrypted response packet len=52
34	8.384563000	192.168.0.11	192.168.0.2	SSHv2	134	Encrypted request packet len=68
37	8.386325000	192.168.0.2	192.168.0.11	SSHv2	150	Encrypted response packet len=84
64	13.700774000	192.168.0.11	192.168.0.2	SSHv2	214	Encrypted request packet len=148
65	13.710428000	192.168.0.2	192.168.0.11	SSHv2	102	Encrypted response packet len=36
67	13.710847000	192.168.0.11	192.168.0.2	SSHv2	186	Encrypted request packet len=120
68	13.714328000	192.168.0.2	192.168.0.11	SSHv2	118	Encrypted response packet len=52
69	13.714663000	192.168.0.11	192.168.0.2	SSHv2	594	Encrypted request packet len=528
70	13.715101000	192.168.0.2	192.168.0.11	SSHv2	174	Encrypted response packet len=108
71	13.716042000	192.168.0.2	192.168.0.11	SSHv2	166	Encrypted response packet len=100
73	13.759743000	192.168.0.2	192.168.0.11	SSHv2	134	Encrypted response packet len=68

Here is the accounting of SSH traffic:


```
[root@DataComm ~]# iptables -L ssh-traffic -v -n -x
Chain ssh-traffic (5 references)
pkts      bytes target    prot opt in      out     source        destination
37         0 ACCEPT    all  --  *      *       0.0.0.0/0     0.0.0.0/0
0          0 ACCEPT    tcp  --  *      *       0.0.0.0/0     192.168.0.2    tcp dpt:22 state NEW,ESTABLISHED
0          0 ACCEPT    tcp  --  *      *       192.168.0.2   0.0.0.0/0     tcp spt:22 state ESTABLISHED
0          0 ACCEPT    tcp  --  *      *       192.168.0.2   0.0.0.0/0     tcp dpt:22 state NEW,ESTABLISHED
0          0 ACCEPT    tcp  --  *      *       0.0.0.0/0     192.168.0.2    tcp spt:22 state ESTABLISHED
[root@DataComm ~]#
```

(5) WWW(80) traffic is allowed and accounted for

From the testing terminal, here are the results of the hping command to our personal firewall:

```
[root@DataComm ~]# hping3 192.168.0.2 -p 80 -c 5
HPING 192.168.0.2 (em1 192.168.0.2): NO FLAGS are set, 40 headers + 0 data bytes
len=46 ip=192.168.0.2 ttl=64 DF id=48618 sport=80 flags=RA seq=0 win=0 rtt=0.3 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48619 sport=80 flags=RA seq=1 win=0 rtt=0.3 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48620 sport=80 flags=RA seq=2 win=0 rtt=0.3 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48621 sport=80 flags=RA seq=3 win=0 rtt=0.4 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48622 sport=80 flags=RA seq=4 win=0 rtt=0.4 ms

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 5 packets received, 0% packet loss
round-trip min/avg/max = 0.3/0.3/0.4 ms
[root@DataComm ~]#
```

Here are the results of the machine containing our firewall sending outbound packets:

```
[root@DataComm ~]# hping3 192.168.0.11 -S -p 80 -c 5
HPING 192.168.0.11 (em1 192.168.0.11): S set, 40 headers + 0 data bytes
len=46 ip=192.168.0.11 ttl=64 DF id=0 sport=80 flags=SA seq=0 win=29200 rtt=0.3 ms
len=46 ip=192.168.0.11 ttl=64 DF id=0 sport=80 flags=SA seq=1 win=29200 rtt=0.3 ms
len=46 ip=192.168.0.11 ttl=64 DF id=0 sport=80 flags=SA seq=2 win=29200 rtt=0.4 ms
len=46 ip=192.168.0.11 ttl=64 DF id=0 sport=80 flags=SA seq=3 win=29200 rtt=0.4 ms
len=46 ip=192.168.0.11 ttl=64 DF id=0 sport=80 flags=SA seq=4 win=29200 rtt=0.4 ms

--- 192.168.0.11 hping statistic ---
5 packets transmitted, 5 packets received, 0% packet loss
round-trip min/avg/max = 0.3/0.4/0.4 ms
[root@DataComm ~]#
```

And here are the accounting information regarding WWW(80) traffic using iptables:

```
[root@DataComm ~]# iptables -L www-traffic -v -x -n
Chain www-traffic (5 references)
pkts      bytes target    prot opt in      out     source        destination
58         0 DROP      all  --  *      *       0.0.0.0/0     0.0.0.0/0
0          0 DROP      tcp  --  em1    *       0.0.0.0/0     192.168.0.2    tcp spts:0:1023 dpt:80
0          0 DROP      tcp  --  em1    *       0.0.0.0/0     192.168.0.2    tcp spts:0:1023 dpt:443
5         200 ACCEPT    tcp  --  *      *       0.0.0.0/0     192.168.0.2    multiport dports 80,443 state NEW,ESTABLISHED
5         200 ACCEPT    tcp  --  *      *       192.168.0.2   0.0.0.0/0     multiport sports 80,443 state ESTABLISHED
24        5172 ACCEPT    tcp  --  *      *       192.168.0.2   0.0.0.0/0     multiport dports 80,443 state NEW,ESTABLISHED
24        1775 ACCEPT    tcp  --  *      *       0.0.0.0/0     192.168.0.2    multiport sports 80,443 state ESTABLISHED
[root@DataComm ~]#
```

(6) WWW(443) traffic is allowed and accounted for

From the testing terminal, here are the results of the hping command to our personal firewall:

```
[root@DataComm ~]# hping3 192.168.0.2 -S -p 443 -c 5
HPING 192.168.0.2 (em1 192.168.0.2): S set, 40 headers + 0 data bytes
len=46 ip=192.168.0.2 ttl=64 DF id=48638 sport=443 flags=RA seq=0 win=0 rtt=0.4 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48639 sport=443 flags=RA seq=1 win=0 rtt=0.4 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48640 sport=443 flags=RA seq=2 win=0 rtt=0.3 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48641 sport=443 flags=RA seq=3 win=0 rtt=0.3 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48642 sport=443 flags=RA seq=4 win=0 rtt=0.3 ms

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 5 packets received, 0% packet loss
round-trip min/avg/max = 0.3/0.4/0.4 ms
[root@DataComm ~]#
```

Here are the results of the machine containing our firewall sending outbound packets:

```
[root@DataComm ~]# hping3 192.168.0.11 -S -p 443 -c 5
HPING 192.168.0.11 (em1 192.168.0.11): S set, 40 headers + 0 data bytes
len=46 ip=192.168.0.11 ttl=64 DF id=18612 sport=443 flags=RA seq=0 win=0 rtt=0.3 ms
len=46 ip=192.168.0.11 ttl=64 DF id=18613 sport=443 flags=RA seq=1 win=0 rtt=0.3 ms
len=46 ip=192.168.0.11 ttl=64 DF id=18614 sport=443 flags=RA seq=2 win=0 rtt=0.3 ms
len=46 ip=192.168.0.11 ttl=64 DF id=18615 sport=443 flags=RA seq=3 win=0 rtt=0.4 ms
len=46 ip=192.168.0.11 ttl=64 DF id=18616 sport=443 flags=RA seq=4 win=0 rtt=0.3 ms

--- 192.168.0.11 hping statistic ---
5 packets transmitted, 5 packets received, 0% packet loss
round-trip min/avg/max = 0.3/0.3/0.4 ms
[root@DataComm ~]#
```

And here are the accounting information regarding WWW(80) traffic using iptables:

```
[root@DataComm ~]# iptables -L www-traffic -v -x -n
Chain www-traffic (5 references)
pkts    bytes target      prot opt in      out     source         destination
268      88768 all -- *          *      0.0.0.0/0      0.0.0.0/0
0         0 DROP tcp -- em1      *      0.0.0.0/0      192.168.0.2    tcp spts:0:1023 dpt:80
0         0 DROP tcp -- em1      *      0.0.0.0/0      192.168.0.2    tcp spts:0:1023 dpt:443
10        400 ACCEPT tcp -- *        *      0.0.0.0/0      192.168.0.2    multiport dports 80,443 state NEW,ESTABLISHED
10        400 ACCEPT tcp -- *        *      192.168.0.2    0.0.0.0/0      multiport sports 80,443 state ESTABLISHED
113      21322 ACCEPT tcp -- *        *      192.168.0.2    0.0.0.0/0      multiport dports 80,443 state NEW,ESTABLISHED
135      66646 ACCEPT tcp -- *        *      0.0.0.0/0      192.168.0.2    multiport sports 80,443 state ESTABLISHED
[root@DataComm ~]#
```

Note the increase of our packet count from 5 to 10.

(7) WWW(80) traffic is disallowed and accounted for when source ports are less than 1024

Here, we set our source port to the left-most extreme of our constraint (port 0). We will execute this on our non-firewalled terminal:

```
[root@DataComm ~]# hping3 192.168.0.2 -S -p 80 -c 5 -s 0
HPING 192.168.0.2 (em1 192.168.0.2): S set, 40 headers + 0 data bytes

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
[root@DataComm ~]#
```

Our packets are dropped here. Let's take a look at our packet auditing on our firewalled terminal:

```
[root@DataComm ~]# iptables -L www-traffic -n -v -x
Chain www-traffic (5 references)
pkts    bytes target     prot opt in     out     source        destination
1599    395621 all -- *      *      0.0.0.0/0     0.0.0.0/0
5        200 DROP      tcp -- em1    *      0.0.0.0/0     192.168.0.2    tcp spts:0:1023 dpt:80
0         0 DROP      tcp -- em1    *      0.0.0.0/0     192.168.0.2    tcp spts:0:1023 dpt:443
10       400 ACCEPT    tcp -- *      *      0.0.0.0/0     192.168.0.2    multiport dports 80,443 state
```

(8) WWW(80) traffic is disallowed and accounted for when source ports are less than 1024

We will redo the test, except we will attempt our right-most extreme (port 1023).

NOTE: we have to specify the --keep switch, otherwise, our ports will increment

```
[root@DataComm ~]# hping3 192.168.0.2 -S -p 80 -c 5 -s 1023 --keep
HPING 192.168.0.2 (em1 192.168.0.2): S set, 40 headers + 0 data bytes

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
[root@DataComm ~]#
```

(9) WWW(80) traffic is disallowed and accounted for when source ports are less than 1024

We will redo our previous test, except this time, we will try a valid port (port 1024+)

NOTE: we won't have to specify --keep switch, because anything greater than 1024 is permitted

```
[root@DataComm ~]# hping3 192.168.0.2 -S -p 80 -c 5 -s 1024
HPING 192.168.0.2 (em1 192.168.0.2): S set, 40 headers + 0 data bytes
len=46 ip=192.168.0.2 ttl=64 DF id=48711 sport=80 flags=RA seq=0 win=0 rtt=0.3 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48712 sport=80 flags=RA seq=1 win=0 rtt=0.3 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48713 sport=80 flags=RA seq=2 win=0 rtt=0.3 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48714 sport=80 flags=RA seq=3 win=0 rtt=0.3 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48715 sport=80 flags=RA seq=4 win=0 rtt=0.4 ms

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 5 packets received, 0% packet loss
round-trip min/avg/max = 0.3/0.3/0.4 ms
[root@DataComm ~]#
```

(10) WWW(443) traffic is disallowed and accounted for when source ports are less than 1024

We will flush our accounting traffic and restart the last 3 tests again for port 443 (https)

Starting with our left-most extreme (port 0):

```
[root@DataComm ~]# hping3 192.168.0.2 -S -p 443 -c 5 -s 0
HPING 192.168.0.2 (em1 192.168.0.2): S set, 40 headers + 0 data bytes

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
[root@DataComm ~]#
```

And then checking our accounting tables:

```
[root@DataComm ~]# iptables -L www-traffic -n -v -x
Chain www-traffic (5 references)

```

pkts	bytes	target	prot	opt	in	out	source	destination	
306	77018		all	--	*	*	0.0.0.0/0	0.0.0.0/0	
0	0	DROP	tcp	--	em1	*	0.0.0.0/0	192.168.0.2	tcp spts:0:1023 dpt:80
5	200	DROP	tcp	--	em1	*	0.0.0.0/0	192.168.0.2	tcp spts:0:1023 dpt:443
0	0	ACCEPT	tcp	--	*	*	0.0.0.0/0	192.168.0.2	multiport dports 80,443

(11) WWW(443) traffic is disallowed and accounted for when source ports are less than 1024

NOTE: we have to specify --keep switch, otherwise our source port will increment

Starting with our right-most extreme (port 1023):

```
[root@DataComm ~]# hping3 192.168.0.2 -S -p 443 -c 5 -s 1023 --keep
HPING 192.168.0.2 (em1 192.168.0.2): S set, 40 headers + 0 data bytes

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
[root@DataComm ~]#
```

(12) WWW(443) traffic is disallowed and accounted for when source ports are less than 1024

NOTE: we won't have to specify --keep switch, because anything greater than 1024 is permitted

Using port 1024:

```
[root@DataComm ~]# hping3 192.168.0.2 -S -p 443 -c 5 -s 1024
HPING 192.168.0.2 (em1 192.168.0.2): S set, 40 headers + 0 data bytes
len=46 ip=192.168.0.2 ttl=64 DF id=48716 sport=443 flags=RA seq=0 win=0 rtt=0.4 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48717 sport=443 flags=RA seq=1 win=0 rtt=0.3 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48718 sport=443 flags=RA seq=2 win=0 rtt=0.4 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48719 sport=443 flags=RA seq=3 win=0 rtt=0.3 ms
len=46 ip=192.168.0.2 ttl=64 DF id=48720 sport=443 flags=RA seq=4 win=0 rtt=0.3 ms

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 5 packets received, 0% packet loss
round-trip min/avg/max = 0.3/0.4/0.4 ms
[root@DataComm ~]#
```

(13) Incoming packets to port 0 are dropped (TCP)

Here is our command:

```
[root@DataComm ~]# hping3 192.168.0.2 -S -p 0 -c 5
HPING 192.168.0.2 (em1 192.168.0.2): S set, 40 headers + 0 data bytes

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
[root@DataComm ~]#
```

As you can see, no responses. Here is our accounting table to prove it:

```
[root@DataComm ~]# iptables -L -n -v -x
Chain INPUT (policy DROP 18 packets, 1867 bytes)
  pkts    bytes target     prot opt in     out     source                 destination            tcp spt:22
    0         0 ssh-traffic tcp -- *      *      0.0.0.0/0             0.0.0.0/0              multiport sp
   12       966 www-traffic tcp -- *      *      0.0.0.0/0             0.0.0.0/0              multiport sp
   84     9198 noness-traffic all -- *      *      0.0.0.0/0             0.0.0.0/0              multiport sp
    5       200 DROP      tcp -- em1    *      0.0.0.0/0             192.168.0.2            tcp dpt:0
    0         0 DROP      udp -- em1    *      0.0.0.0/0             192.168.0.2            udp dpt:0
    0         0 DROP      tcp -- em1    *      0.0.0.0/0             0.0.0.0/0              tcp flags:!!0x1
    0         0 www-traffic tcp -- em1    *      0.0.0.0/0             0.0.0.0/0              multiport dp
    0         0 www-traffic tcp -- em1    *      0.0.0.0/0             0.0.0.0/0              multiport sp
   59     6697 ssh-traffic tcp -- em1    *      0.0.0.0/0             0.0.0.0/0              tcp dpt:22
    0         0 ssh-traffic tcp -- em1    *      0.0.0.0/0             0.0.0.0/0              tcp spt:22
    0         0 noness-traffic tcp -- em1    *      0.0.0.0/0             0.0.0.0/0              multiport
    0         0 noness-traffic tcp -- em1    *      0.0.0.0/0             0.0.0.0/0              multiport
   18     1867 noness-traffic udp -- em1    *      0.0.0.0/0             0.0.0.0/0              multiport
   18     1867 noness-traffic udp -- em1    *      0.0.0.0/0             0.0.0.0/0              multiport
```

(14) Incoming packets to port 0 are dropped (UDP)

Here is our command; note the --udp switch:

```
[root@DataComm ~]# hping3 192.168.0.2 --udp -p 0 -c 5
HPING 192.168.0.2 (em1 192.168.0.2): udp mode set, 28 headers + 0 data bytes

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
[root@DataComm ~]#
```

As you can see, no responses. Here is our accounting table to prove it; note the difference between the TCP and UDP count:

```
[root@DataComm ~]# iptables -L -n -v -x
```

Chain INPUT (policy DROP 69 packets, 8629 bytes)									
pkts	bytes	target	prot	opt	in	out	source	destination	
0	0	ssh-traffic	tcp	--	*	*	0.0.0.0/0	0.0.0.0/0	tcp spt:22
430	176693	www-traffic	tcp	--	*	*	0.0.0.0/0	0.0.0.0/0	multiport s
242	27184	noness-traffic	all	--	*	*	0.0.0.0/0	0.0.0.0/0	
5	200	DROP	tcp	--	em1	*	0.0.0.0/0	192.168.0.2	tcp dpt:0
5	140	DROP	udp	--	em1	*	0.0.0.0/0	192.168.0.2	udp dpt:0
0	0	DROP	tcp	--	em1	*	0.0.0.0/0	0.0.0.0/0	tcp flags: !0x
0	0	www-traffic	tcp	--	em1	*	0.0.0.0/0	0.0.0.0/0	multiport d
0	0	www-traffic	tcp	--	em1	*	0.0.0.0/0	0.0.0.0/0	multiport s
153	15830	ssh-traffic	tcp	--	em1	*	0.0.0.0/0	0.0.0.0/0	tcp dpt:22
0	0	ssh-traffic	tcp	--	em1	*	0.0.0.0/0	0.0.0.0/0	tcp spt:22
0	0	noness-traffic	tcp	--	em1	*	0.0.0.0/0	0.0.0.0/0	multiport
0	0	noness-traffic	tcp	--	em1	*	0.0.0.0/0	0.0.0.0/0	multiport
69	8629	noness-traffic	udp	--	em1	*	0.0.0.0/0	0.0.0.0/0	multiport
69	8629	noness-traffic	udp	--	em1	*	0.0.0.0/0	0.0.0.0/0	multiport

(15) WWW(80) incoming traffic is accounted for

In our non-firewall terminal, we ran this command:

hping3 192.168.0.2 -S -p 80 -c 5

Here is our accounting table:

```
[root@DataComm ~]# iptables -L www-traffic -n -v -x
```

Chain www-traffic (5 references)									
pkts	bytes	target	prot	opt	in	out	source	destination	
83	14624	all	--	*	*	*	0.0.0.0/0	0.0.0.0/0	
0	0	DROP	tcp	--	em1	*	0.0.0.0/0	192.168.0.2	tcp spts:0:1023 dpt:80
0	0	DROP	tcp	--	em1	*	0.0.0.0/0	192.168.0.2	tcp spts:0:1023 dpt:443
5	200	ACCEPT	tcp	--	*	*	0.0.0.0/0	192.168.0.2	multiport dports 80,443 state NEW,ESTABLISHED
5	200	ACCEPT	tcp	--	*	*	192.168.0.2	0.0.0.0/0	multiport sports 80,443 state ESTABLISHED
39	7078	ACCEPT	tcp	--	*	*	192.168.0.2	0.0.0.0/0	multiport dports 80,443 state NEW,ESTABLISHED
34	7146	ACCEPT	tcp	--	*	*	0.0.0.0/0	192.168.0.2	multiport sports 80,443 state ESTABLISHED

(16) WWW(80) outgoing traffic is accounted for

In our firewallled terminal, we ran this command:

hping3 192.168.0.11 -S -p 80 -c 5

Here is our accounting table:

```
[root@DataComm ~]# iptables -L www-traffic -n -v -x
```

Chain www-traffic (5 references)									
pkts	bytes	target	prot	opt	in	out	source	destination	
83	14624	all	--	*	*	*	0.0.0.0/0	0.0.0.0/0	
0	0	DROP	tcp	--	em1	*	0.0.0.0/0	192.168.0.2	tcp spts:0:1023 dpt:80
0	0	DROP	tcp	--	em1	*	0.0.0.0/0	192.168.0.2	tcp spts:0:1023 dpt:443
5	200	ACCEPT	tcp	--	*	*	0.0.0.0/0	192.168.0.2	multiport dports 80,443 state NEW,ESTABLISHED
5	200	ACCEPT	tcp	--	*	*	192.168.0.2	0.0.0.0/0	multiport sports 80,443 state ESTABLISHED
39	7078	ACCEPT	tcp	--	*	*	192.168.0.2	0.0.0.0/0	multiport dports 80,443 state NEW,ESTABLISHED
34	7146	ACCEPT	tcp	--	*	*	0.0.0.0/0	192.168.0.2	multiport sports 80,443 state ESTABLISHED

NOTE the two different highlighted lines: the former depicts incoming while the latter depicts outgoing.

(17) WWW(443) incoming traffic is accounted for
 In our non-firewall terminal, we ran this command:
 hping3 192.168.0.2 -S -p 443 -c 5
 Here is our accounting table:

```
[root@DataComm ~]# iptables -L www-traffic -n -v -x
```

Chain www-traffic (5 references)									
pkts	bytes	target	prot	opt	in	out	source	destination	
1781	439307	all	--	*	*	*	0.0.0.0/0	0.0.0.0/0	
0	0	DROP	tcp	--	em1	*	0.0.0.0/0	192.168.0.2	tcp spts:0:1023 dpt:80
0	0	DROP	tcp	--	em1	*	0.0.0.0/0	192.168.0.2	tcp spts:0:1023 dpt:443
10	400	ACCEPT	tcp	--	*	*	0.0.0.0/0	192.168.0.2	multiport dports 80,443 state NEW,ESTABLISHED
10	400	ACCEPT	tcp	--	*	*	192.168.0.2	0.0.0.0/0	multiport sports 80,443 state ESTABLISHED
855	83242	ACCEPT	tcp	--	*	*	192.168.0.2	0.0.0.0/0	multiport dports 80,443 state NEW,ESTABLISHED
906	355265	ACCEPT	tcp	--	*	*	0.0.0.0/0	192.168.0.2	multiport sports 80,443 state ESTABLISHED

```
[root@DataComm ~]#
```

(18) WWW(443) outgoing traffic is accounted for
 In our firewalled terminal, we ran this command:
 hping3 192.168.0.11 -S -p 443 -c 5
 Here is our accounting table:

```
[root@DataComm ~]# iptables -L www-traffic -n -v -x
```

Chain www-traffic (5 references)									
pkts	bytes	target	prot	opt	in	out	source	destination	
1781	439307	all	--	*	*	*	0.0.0.0/0	0.0.0.0/0	
0	0	DROP	tcp	--	em1	*	0.0.0.0/0	192.168.0.2	tcp spts:0:1023 dpt:80
0	0	DROP	tcp	--	em1	*	0.0.0.0/0	192.168.0.2	tcp spts:0:1023 dpt:443
10	400	ACCEPT	tcp	--	*	*	0.0.0.0/0	192.168.0.2	multiport dports 80,443 state NEW,ESTABLISHED
10	400	ACCEPT	tcp	--	*	*	192.168.0.2	0.0.0.0/0	multiport sports 80,443 state ESTABLISHED
855	83242	ACCEPT	tcp	--	*	*	192.168.0.2	0.0.0.0/0	multiport dports 80,443 state NEW,ESTABLISHED
906	355265	ACCEPT	tcp	--	*	*	0.0.0.0/0	192.168.0.2	multiport sports 80,443 state ESTABLISHED

```
[root@DataComm ~]#
```

NOTE the two different highlighted lines: the former depicts incoming while the latter depicts outgoing. Also note the increase of packets from 5 to 10.

(19) SSH(22) incoming traffic is accounted for
 In our non-firewalled terminal, we ran this command:
 hping3 192.168.0.2 -S -p 22 -c 5
 Here is our accounting table:

```
[root@DataComm ~]# iptables -L ssh-traffic -n -v -x
```

Chain ssh-traffic (5 references)									
pkts	bytes	target	prot	opt	in	out	source	destination	
30	1240	all	--	*	*	*	0.0.0.0/0	0.0.0.0/0	
10	400	ACCEPT	tcp	--	*	*	0.0.0.0/0	192.168.0.2	tcp dpt:22 state NEW,ESTABLISHED
5	220	ACCEPT	tcp	--	*	*	192.168.0.2	0.0.0.0/0	tcp spt:22 state ESTABLISHED
10	400	ACCEPT	tcp	--	*	*	192.168.0.2	0.0.0.0/0	tcp dpt:22 state NEW,ESTABLISHED
5	220	ACCEPT	tcp	--	*	*	0.0.0.0/0	192.168.0.2	tcp spt:22 state ESTABLISHED

```
[root@DataComm ~]#
```

(20) SSH(22) outgoing traffic is accounted for

In our firewalled terminal, we ran this command:

```
hping3 192.168.0.11 -S -p 22 -c 5
```

Here is our accounting table:

```
[root@DataComm ~]# iptables -L ssh-traffic -n -v -x
Chain ssh-traffic (5 references)
pkts    bytes target     prot opt in     out     source           destination
 30      1240      all -- *      *      0.0.0.0/0        0.0.0.0/0
 10       400 ACCEPT tcp -- *      *      0.0.0.0/0        192.168.0.2      tcp dpt:22 state NEW,ESTABLISHED
 5        220 ACCEPT tcp -- *      *      192.168.0.2      0.0.0.0/0        tcp spt:22 state ESTABLISHED
 10       400 ACCEPT tcp -- *      *      192.168.0.2      0.0.0.0/0        tcp dpt:22 state NEW,ESTABLISHED
 5        220 ACCEPT tcp -- *      *      0.0.0.0/0        192.168.0.2      tcp spt:22 state ESTABLISHED
[root@DataComm ~]#
```

NOTE the two different highlighted lines: the former depicts incoming while the latter depicts outgoing.

(21) All incoming SYN packets are dropped unless permitted

In our non-firewalled terminal, we ran this command with the following results:

```
[root@DataComm ~]# hping3 192.168.0.2 -S -c 5
HPING 192.168.0.2 (em1 192.168.0.2): S set, 40 headers + 0 data bytes

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
[root@DataComm ~]#
```

NOTE: There is no -p switch. We assume that SYN packets that are allowed are specified in previous test cases. Any other unspecified TCP packets coming in are therefore dropped because they are not explicitly permitted.

(22) All other undefined ports with inbound traffic are accounted for (TCP)

We ran hping3 with an undefined port of 8080:

```
[root@DataComm ~]# hping3 192.168.0.2 -S -p 8080 -c 5
HPING 192.168.0.2 (em1 192.168.0.2): S set, 40 headers + 0 data bytes

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
```

We expect that it will drop these packets but we will see that it is accounted for in our tables:


```
[root@DataComm ~]# iptables -L INPUT -n -v -x
Chain INPUT (policy DROP 11 packets, 1131 bytes)
pkts    bytes target     prot opt in     out     source            destination
0        0 ssh-traffic tcp -- *      *      0.0.0.0/0         0.0.0.0/0         tcp spt:22
31      6745 www-traffic tcp -- *      *      0.0.0.0/0         0.0.0.0/0         multiport sports 80,443
12     1368 noness-traffic all -- *      *      0.0.0.0/0         0.0.0.0/0
0        0 DROP      tcp -- em1    *      0.0.0.0/0         192.168.0.2        tcp dpt:0
0        0 DROP      udp -- em1    *      0.0.0.0/0         192.168.0.2        udp dpt:0
0        0 DROP      tcp -- em1    *      0.0.0.0/0         0.0.0.0/0         tcp flags:!0x17/0x02 state NEW
0        0 www-traffic tcp -- em1    *      0.0.0.0/0         0.0.0.0/0         multiport dports 80,443
0        0 www-traffic tcp -- em1    *      0.0.0.0/0         0.0.0.0/0         multiport sports 80,443
0        0 ssh-traffic tcp -- em1    *      0.0.0.0/0         0.0.0.0/0         tcp dpt:22
0        0 ssh-traffic tcp -- em1    *      0.0.0.0/0         0.0.0.0/0         tcp spt:22
0        0 noness-traffic tcp -- em1    *      192.168.0.2       0.0.0.0/0         multiport dports !80,443,22
5       200 noness-traffic tcp -- em1    *      0.0.0.0/0         192.168.0.2        multiport sports !80,443,22
0        0 noness-traffic udp -- em1    *      192.168.0.2       0.0.0.0/0         multiport dports !80,443,22
0        0 noness-traffic udp -- em1    *      0.0.0.0/0         192.168.0.2        multiport sports !80,443,22
[root@DataComm ~]#
```

(23) All other undefined ports with outbound traffic are accounted for (TCP)

We ran hping3 with an undefined port of 8080 from our firewalled terminal:

```
[root@DataComm ~]# hping3 192.168.0.11 -S -p 8080 -c 5
HPING 192.168.0.11 (em1 192.168.0.11): S set, 40 headers + 0 data bytes
[send_ip] sendto: Operation not permitted
```

We expect that it will drop these packets because we should be a “good Internet neighbour”.

Therefore, our firewall restricts us from being a bad neighbour. However, at least one instance of our “probing” must go through. It is recorded in our tables here:

```
[root@DataComm ~]# iptables -L OUTPUT -n -v -x
Chain OUTPUT (policy DROP 23 packets, 7256 bytes)
pkts    bytes target     prot opt in     out     source            destination
0        0 DROP      tcp -- *      em1    192.168.0.2       0.0.0.0/0         tcp spt:0
0        0 DROP      udp -- *      em1    192.168.0.2       0.0.0.0/0         udp spt:0
0        0 www-traffic tcp -- *      em1    0.0.0.0/0         0.0.0.0/0         multiport sports 80,443
318    173941 www-traffic tcp -- *      em1    0.0.0.0/0         0.0.0.0/0         multiport dports 80,443
0        0 ssh-traffic tcp -- *      em1    0.0.0.0/0         0.0.0.0/0         tcp spt:22
0        0 ssh-traffic tcp -- *      em1    0.0.0.0/0         0.0.0.0/0         tcp dpt:22
1       40 noness-traffic tcp -- *      em1    192.168.0.2       0.0.0.0/0         multiport sports !80,443,22
0        0 noness-traffic tcp -- *      em1    0.0.0.0/0         192.168.0.2        multiport dports !80,443,22
26     7474 noness-traffic udp -- *      em1    192.168.0.2       0.0.0.0/0         multiport sports !80,443,22
0        0 noness-traffic udp -- *      em1    0.0.0.0/0         192.168.0.2        multiport dports !80,443,22
[root@DataComm ~]#
```

Note the one instance of outbound traffic.

(24) All other undefined ports with inbound traffic are accounted for (UDP)

We ran hping3 with an undefined port of 8080 with the --udp switch from our non-firewalled terminal:

```
[root@DataComm ~]# hping3 192.168.0.2 --udp -p 8080 -c 5
HPING 192.168.0.2 (em1 192.168.0.2): udp mode set, 28 headers + 0 data bytes

--- 192.168.0.2 hping statistic ---
5 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
[root@DataComm ~]#
```

We expect that it will drop these packets but we will see that it is accounted for in our tables:

```
[root@DataComm ~]# iptables -L INPUT -n -v -x
Chain INPUT (policy DROP 143 packets, 18807 bytes)
pkts    bytes target      prot opt in     out     source            destination        tcp spt:22
2681    471749 www-traffic  tcp  --  *      *      0.0.0.0/0         0.0.0.0/0          multiport sports 80,443
158     22691 noness-traffic all  --  *      *      0.0.0.0/0         0.0.0.0/0          tcp dpt:0
0       0 DROP      tcp  --  em1    *      0.0.0.0/0         192.168.0.2        tcp dpt:0
0       0 DROP      udp  --  em1    *      0.0.0.0/0         192.168.0.2        udp dpt:0
0       0 DROP      tcp  --  em1    *      0.0.0.0/0         0.0.0.0/0          tcp flags:!0x17/0x02 state NEW
0       0 www-traffic tcp  --  em1    *      0.0.0.0/0         0.0.0.0/0          multiport dports 80,443
1       64 www-traffic tcp  --  em1    *      0.0.0.0/0         0.0.0.0/0          multiport sports 80,443
0       0 ssh-traffic tcp  --  em1    *      0.0.0.0/0         0.0.0.0/0          tcp dpt:22
0       0 ssh-traffic tcp  --  em1    *      0.0.0.0/0         0.0.0.0/0          tcp dpt:22
0       0 noness-traffic tcp -- em1    *      192.168.0.2       0.0.0.0/0          multiport dports !80,443,22
5       200 noness-traffic tcp -- em1    *      0.0.0.0/0         192.168.0.2        multiport sports !80,443,22
0       0 noness-traffic udp  --  em1    *      192.168.0.2       0.0.0.0/0          multiport dports !80,443,22
5       140 noness-traffic udp  --  em1    *      0.0.0.0/0         192.168.0.2        multiport sports !80,443,22
[root@DataComm ~]#
```

(25) All other undefined ports with outbound traffic are accounted for (UDP)

We ran hping3 with an undefined port of 8080 with the --udp switch from our firewalled terminal:

```
[root@DataComm ~]# hping3 192.168.0.11 --udp -p 8080 -c 5
HPING 192.168.0.11 (em1 192.168.0.11): udp mode set, 28 headers + 0 data bytes
[send_ip] sendto: Operation not permitted
[root@DataComm ~]#
```

We expect this to happen because we are not being a “good Internet neighbour”. Thus, our firewall blocks our attempt to probe. However, because our UDP rules are universal, we will have to check a before and after screenshot of our tables:

```
[root@DataComm ~]# iptables -L OUTPUT -n -v -x
Chain OUTPUT (policy DROP 51 packets, 15336 bytes)
pkts    bytes target      prot opt in     out     source            destination        tcp spt:0
0       0 DROP      tcp  --  *      em1    192.168.0.2       0.0.0.0/0          tcp spt:0
0       0 DROP      udp  --  *      em1    192.168.0.2       0.0.0.0/0          udp spt:0
0       0 www-traffic tcp  --  *      em1    0.0.0.0/0         0.0.0.0/0          multiport sports 80,443
2655    1162079 www-traffic tcp  --  *      em1    0.0.0.0/0         0.0.0.0/0          multiport dports 80,443
0       0 ssh-traffic tcp  --  *      em1    0.0.0.0/0         0.0.0.0/0          tcp spt:22
0       0 ssh-traffic tcp  --  *      em1    0.0.0.0/0         0.0.0.0/0          tcp dpt:22
1       40 noness-traffic tcp -- *      em1    192.168.0.2       0.0.0.0/0          multiport sports !80,443,22
0       0 noness-traffic tcp -- *      em1    0.0.0.0/0         192.168.0.2        multiport dports !80,443,22
60      15972 noness-traffic udp  --  *      em1    192.168.0.2       0.0.0.0/0          multiport sports !80,443,22
0       0 noness-traffic udp  --  *      em1    0.0.0.0/0         192.168.0.2        multiport dports !80,443,22
[root@DataComm ~]# hping3 192.168.0.11 --udp -p 8080 -c 5
```

Note the 60...

After running another instance of our hping3 command with the same switches, and then re-running our command to display our accounting tables, note the increment in outbound UDP packets:

```
[root@DataComm ~]# hping3 192.168.0.11 --udp -p 8080 -c 5
HPING 192.168.0.11 (em1 192.168.0.11): udp mode set, 28 headers + 0 data bytes
[send_ip] sendto: Operation not permitted
[root@DataComm ~]# iptables -L OUTPUT -n -v -x
Chain OUTPUT (policy DROP 52 packets, 15364 bytes)
pkts    bytes target      prot opt in     out     source            destination        tcp spt:0
0       0 DROP      tcp  --  *      em1    192.168.0.2       0.0.0.0/0          tcp spt:0
0       0 DROP      udp  --  *      em1    192.168.0.2       0.0.0.0/0          udp spt:0
0       0 www-traffic tcp  --  *      em1    0.0.0.0/0         0.0.0.0/0          multiport sports 80,443
2662    1164668 www-traffic tcp  --  *      em1    0.0.0.0/0         0.0.0.0/0          multiport dports 80,443
0       0 ssh-traffic tcp  --  *      em1    0.0.0.0/0         0.0.0.0/0          tcp spt:22
0       0 ssh-traffic tcp  --  *      em1    0.0.0.0/0         0.0.0.0/0          tcp dpt:22
1       40 noness-traffic tcp -- *      em1    192.168.0.2       0.0.0.0/0          multiport sports !80,443,22
0       0 noness-traffic tcp -- *      em1    0.0.0.0/0         192.168.0.2        multiport dports !80,443,22
61      16000 noness-traffic udp  --  *      em1    192.168.0.2       0.0.0.0/0          multiport sports !80,443,22
0       0 noness-traffic udp  --  *      em1    0.0.0.0/0         192.168.0.2        multiport dports !80,443,22
[root@DataComm ~]#
```

Similar to our TCP instance, the firewall must capture at least one before hping3 gets terminated after one instance of probing.

Conclusion

Through the extensive test cases of this assignment, it is sufficient to say that our firewall implementation has covered the relevant criteria of Assignment 1. They are the following:

- Set the default policies to DROP.
- Permit inbound and outbound ssh packets.
- Permit inbound and outbound www packets.
- Drop inbound traffic to port 80 (http) from source ports less than 1024.
- Drop all incoming packets from reserved port 0 as well as outbound traffic to port 0.
- Create a set of user-defined chains that will implement accounting rules to keep track of www, ssh traffic, versus the rest of the traffic on your system.
- Use Netfilter for your firewall implementation.
- Ensure the the firewall drops all inbound SYN packets, unless there is a rule that permits inbound traffic.
- Remember to allow DNS and DHCP traffic

We have also included the functionality to check WWW port 443 (https) in our rule set as well.

Appendix

Located on disk are the following:

- Personal Firewalls - An Introduction to Firewall Administration (.pdf)
- Personal Firewalls - Design and Preliminary Testing (.pdf)
- firewall_1.sh
- README.txt